

SC1015

Mini Project

B137 Team 4

Puah Rong Qi (U2221477H)

Kiersten Yeo Shu Xian (U2220969G)

Manasarow Gunasekaran (U2222761C)

MOTIVATION

01

PROBLEM

- ❑ When issuing credit cards, the personal information of applicants is used to decide whether they should be issued a credit card
- ❑ One important factor is the applicant's loan payment history



Which variable affects
a person's ability to
pay off loans on time
the most?

DATASET

application_record.csv (54.34 MB)


Detail Compact Column

About this file

There are two tables, which are connected by **ID**.

Detailed explanation is [here](#)

application_record.csv contains appliers personal information, which y features for predicting.


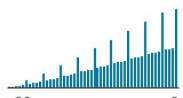
ID	CODE_GENDER	FLAG_OWN_CAR
client number	gender	if users have a car
	F 67%	tru
	M 33%	0
5.01m 8.00m		fal
		0

credit_record.csv (15.37 MB)

Detail Compact Column

About this file


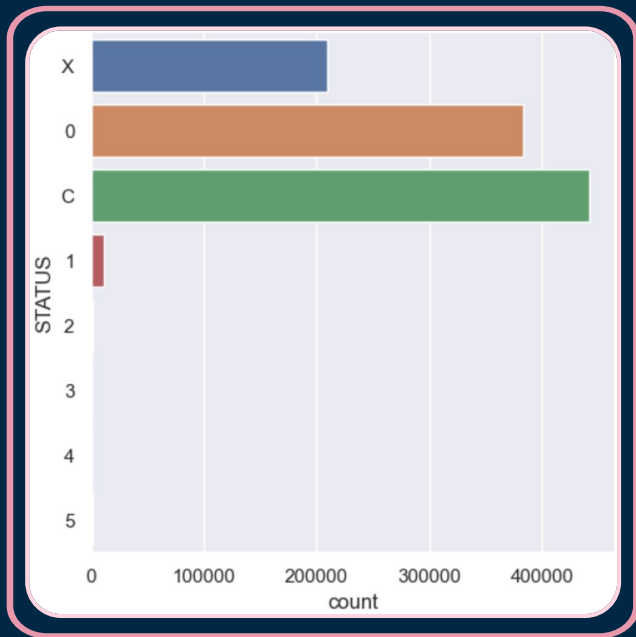
Detailed explanation is [here](#)

ID	# MONTHS_BALAN...	STATUS
ID	record month: The month of the extracted data is the starting point, backwards, 0 is the current month, -1 is the	0: 1-29 days past due 1: 30-59 days past due 2: 60-89 days overdue 3: 90-119 days overdue 4: 120-149 days overdue 5:
		C 42%
5.00m 5.15m	-60 0	O 37%
		Other (223424) 21%

DATA PREPARATION

02

DATA PREPARATION



Changed all the values of
the variable 'STATUS' to numbers

DATA PREPARATION

NULL values

Filled up NULL values
in 'OCCUPATION_TYPE'

```
# Filling the NULL values in Occupation Type in ar  
ar['OCCUPATION_TYPE'].fillna(value="NA", inplace=True)
```

'ID'

Combined rows
with the same IDs

```
# Combining all the repeated IDs and averaging out STATUS  
cr = cr.groupby(['ID'])['STATUS'].agg('mean')
```

Two CSV files

Merged the two dataframes,
removed IDs not in both files

```
# Merging the 2 dataframes and removing all different IDs  
merged_df = pd.merge(cr, ar, on="ID", how="inner")  
# Rounding up STATUS to whole number  
merged_df["STATUS"] = np.ceil(merged_df["STATUS"])
```


EXPLORATORY DATA ANALYSIS

03

VARIABLE SELECTION 1

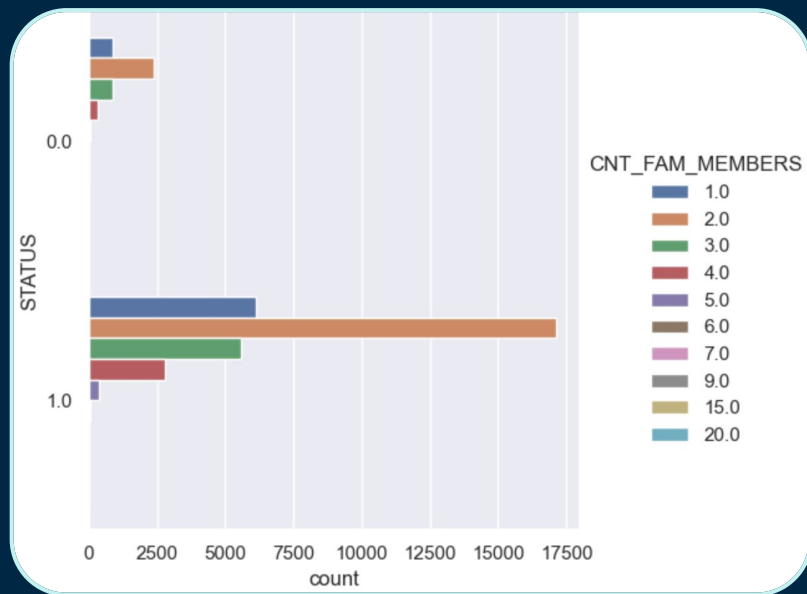
KEEP

- ☐ Total income
- ☐ Education type
- ☐ Marital status
- ☐ Employment status
- ☐ Occupation type
- ☐ Number of family members
- ☐ Gender
- ☐ Number of children
- ☐ Housing type
- ☐ Income type

REMOVE

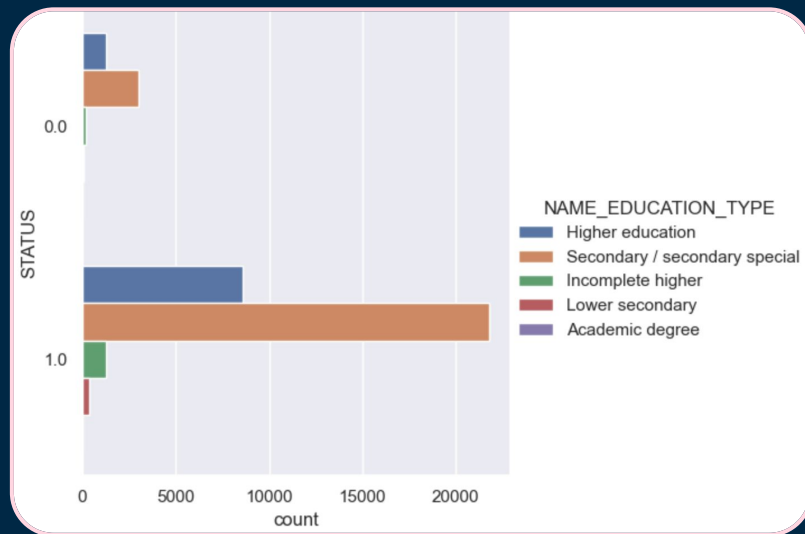
- ☐ Car ownership
- ☐ Property ownership
- ☐ Birthday
- ☐ Phone ownership
- ☐ Email ownership
- ☐ Mobile phone ownership
- ☐ Work phone ownership

CATEGORICAL PLOT



Number of
Family Members

Education Type



CHI-SQUARED TEST

```
create a contingency table of counts
cont_table = pd.crosstab(merged_df["STATUS"], merged_df["OCCUPATION_TYPE"])

# perform chi-square test
chi2, p_val, dof, exp_freq = stats.chi2_contingency(cont_table)

# print results
print("Chi-square test results:")
print(f"Chi-square statistic: {chi2:.2f}")
print(f"P-value: {p_val:.4f}")
print(f"Degrees of freedom: {dof}")
print("Expected frequencies:")
print(exp_freq)
```

```
Chi-square test results:
Chi-square statistic: 41.85
P-value: 0.0012
Degrees of freedom: 18
Expected frequencies:
[[1.51648655e+02  6.73315138e+01  8.00401843e+01  4.38815728e+02
  2.61260938e+02  1.03868941e+01  1.69000878e+02  7.33192528e+00
  7.58976465e+02  2.13847821e+01  3.68062649e+02  1.47493897e+02
  1.38365650e+03  4.20363716e+01  9.65370162e+00  4.25862660e+02
  1.84520120e+01  7.23416628e+01  2.12625833e+01]
[1.08935135e+03  4.83668486e+02  5.74959816e+02  3.15218427e+03
  1.87673906e+03  7.46131059e+01  1.21399912e+03  5.26680747e+01
  5.45202353e+03  1.53615218e+02  2.64393735e+03  1.05950610e+03
  9.93934350e+03  3.01963628e+02  6.93462984e+01  3.05913734e+03
  1.32547988e+02  5.19658337e+02  1.52737417e+02]]
```

Occupation Type

Total Income

```
# create a contingency table of counts
cont_table = pd.crosstab(merged_df["STATUS"], merged_df["Income Range"])

# perform chi-square test
chi2, p_val, dof, exp_freq = stats.chi2_contingency(cont_table)

# print results
print("Chi-square test results:")
print(f"Chi-square statistic: {chi2:.2f}")
print(f"P-value: {p_val:.4f}")
print(f"Degrees of freedom: {dof}")
print("Expected frequencies:")
print(exp_freq)
```

```
Chi-square test results:
Chi-square statistic: 19.48
P-value: 0.0016
Degrees of freedom: 5
Expected frequencies:
[[ 38.84111806  576.56675248 1238.55976462  926.98381758  831.39352703
  335.65502023]
 [ 282.15888194 4188.43324752 8997.44023538 6734.01618242 6039.60647297
 2438.34497977]]
```

VARIABLE SELECTION 2

KEEP

- ☐ Total income
- ☐ Occupation type
- ☐ Education type
- ☐ Number of family members

REMOVE

- ☐ Gender
- ☐ Car ownership
- ☐ Property ownership
- ☐ Number of children
- ☐ Income type
- ☐ Marital status
- ☐ Housing type
- ☐ Birthday
- ☐ Employment status
- ☐ Phone ownership
- ☐ Email ownership
- ☐ Mobile phone ownership
- ☐ Work phone ownership

MACHINE LEARNING

04

TARGET ENCODING

Used to convert categorical variables to numerical variables

```
# Use target encoder to encode the non-numerical columns
# Define the target encoder with smoothing to avoid overfitting
encoder = ce.TargetEncoder(cols=["NAME_INCOME_TYPE", "NAME_EDUCATION_TYPE", "NAME_FAMILY_STATUS", "NAME_HOUSING_TYPE",
                                "OCCUPATION_TYPE", "CODE_GENDER", "FLAG_OWN_CAR", "FLAG_OWN_REALTY"], smoothing=0.2)
print("Converted non-numerical columns to numerical columns using target encoder with smoothing")

# Fit and transform the encoder on the data
merged_df = encoder.fit_transform(merged_df, merged_df["STATUS"])

# Show number of numerical and non-numerical columns in the dataframe
print("\n[POST CONVERSION]\n Number of numerical columns: {}".format(merged_df.select_dtypes(include=np.number).shape[1]))
print("\n[POST CONVERSION]\n Number of non-numerical columns: {}".format(merged_df.select_dtypes(exclude=np.number).shape[1]))
```

Converted non-numerical columns to numerical columns using target encoder with smoothing

```
[POST CONVERSION]
Number of numerical columns: 19
```

```
[POST CONVERSION]
Number of non-numerical columns: 1
```

10

-



UPSCALING

- Initially had a high False Positive rate due to the data for SAMPLE being skewed
- To balance it, we upsampled the data

```
# Upscale Bad to match Good
from sklearn.utils import resample

creditBad = merged_df[merged_df.STATUS == 1]
creditGood = merged_df[merged_df.STATUS == 0]

# Upscale the good samples
creditgood_up = resample(creditGood,
                        replace=True,
                        n_samples=creditBad.shape[0])


# Combine the two classes back after upsampling
scaled_df = pd.concat([creditBad, creditgood_up])

# Check the ratio of the classes
scaled_df['STATUS'].value_counts()


1.0    32002
0.0    32002
Name: STATUS, dtype: int64
```

UPSCALING

Test Data
Accuracy : 0.8773084658987018
TPR Test : 0.9976024184301053
TNR Test : 0.019330855018587362
FPR Test : 0.9806691449814127
FNR Test : 0.002397581569894715




Test Data
Accuracy : 0.5459327153421518
TPR Test : 0.7443270940081564
TNR Test : 0.3491026040045648
FPR Test : 0.6508973959954352
FNR Test : 0.25567290599184356



RANDOM FOREST

- ❑ Builds several decision trees using different samples from the data
- ❑ Takes the majority for classification, average for regression



GRIDSEARCH CROSS-VALIDATION

- ❑ Used to find the optimal values for hyperparameters
- ❑ Using hyperparameters best suited to the model helps to increase the model's accuracy

```
GridSearchCV(cv=5, estimator=RandomForestClassifier(),  
             param_grid={'max_depth': array([14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24]),  
                         'n_estimators': array([100, 200, 300, 400, 500])},  
             scoring='balanced_accuracy')
```

```
RandomForestClassifier(max_depth=21, n_estimators=500)  
0.5983409389822494
```

INSIGHTS

05

The variable that affects a
person's ability to pay off
loans on time the most is
TOTAL INCOME.

NEW TOOLS AND TECHNIQUES



`.MERGE`

Used to merge two
dataframes
together



UPSCALE

Used to balance
skewed data



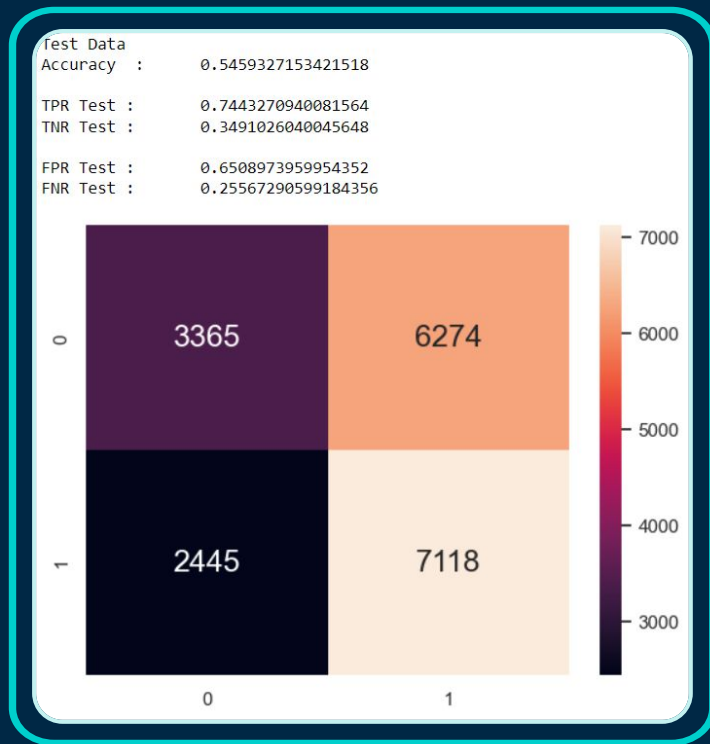
MACHINE LEARNING

Random Forest,
GridSearchCV

Best model for our problem:

Random Forest with
GridSearchCV

RECOMMENDATIONS



DECISION TREE



RANDOM FOREST

RECOMMENDATIONS

DECISION TREE

Single model

Searches for the locally optimal solution at each step

**RANDOM
FOREST**

Multiple models

More likely to find the globally optimal solution

RECOMMENDATIONS



USE MORE VARIABLES

Collecting data on other variables (eg: financial indicators) could improve accuracy



OTHER MACHINE LEARNING MODELS

Advanced models (eg: neural networks) may be better suited to solve this problem

The background is a dark navy blue. It is decorated with various geometric elements: small squares in teal, orange, and pink, some of which are solid and others are hollow outlines. Thin white vertical lines of varying lengths are scattered across the frame. The text 'THANK YOU!' is centered in the middle of the image.

THANK
YOU!

REFERENCES

- ❑ Database:
https://www.kaggle.com/datasets/rikdifos/credit-card-approval-prediction?select=credit_record.csv
- ❑ Decision Tree:
<https://www.ibm.com/docs/en/spss-statistics/25.0.0?topic=trees-creating-decision>
- ❑ Random Forest:
<https://www.analyticsvidhya.com/blog/2021/06/understanding-random-forest/>
- ❑ GridSearchCV: <https://www.mygreatlearning.com/blog/gridsearchcv/>
- ❑ Comparing Decision Tree and Random Forest:
<https://vitalflux.com/differences-between-decision-tree-random-forest/#:~:text=Random%20forest%20is%20a%20ensemble,series%20of%20if%2Dthen%20rules.>